① pointers to pointers on CUDA device is hard to set up and very slow in performance.

— Slower than linearized 1D storage

Therefore do NOT use them unless there is no other option for you!

— We will learn why it is slow later!

② Demo 1: pptr.cu

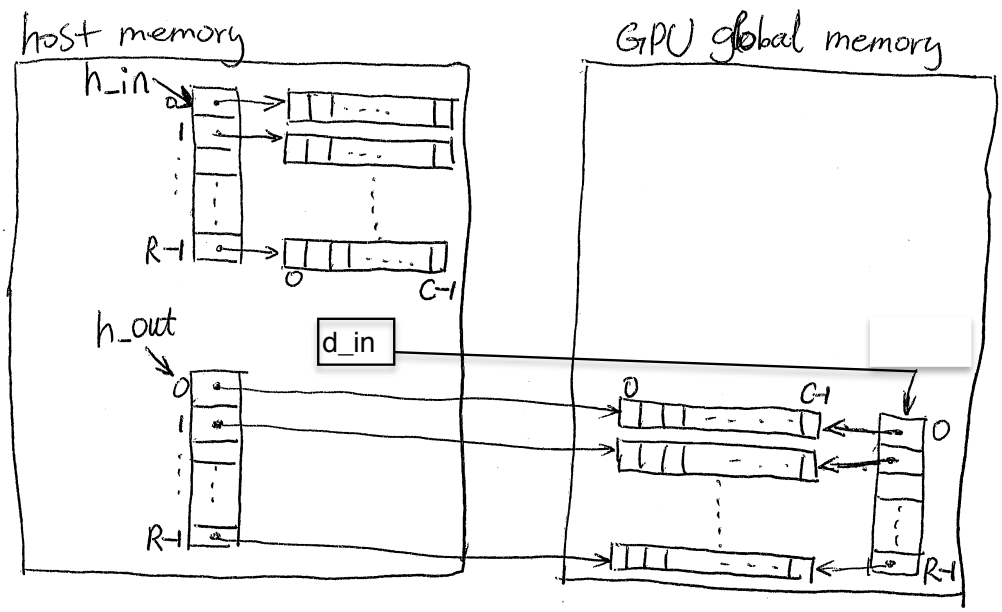

- each time, cudaMemcpy can only copy a piece of <u>contiguous</u> memory from device to host, and vice versa. And the cudaMemcpy needs the start address of that piece of memory and how many bytes to be copied!
- Note that; each row in 2D array (represented by this double pointer) could be far away from each other in memory address.